

Evolving Algebras 1993: Lipari Guide*

Yuri Gurevich[†]

Contents

1	Introduction	3
2	Static Algebras and Updates	4
2.1	Static Algebras: Motivation	4
2.2	Vocabularies	5
2.3	Definition of Static Algebras	5
2.4	Terms	5
2.5	Locations and Updates	6
2.6	Update Sets and Families of Update Sets	6
2.7	Conservative Determinism vs. Local Nondeterminism	7
3	Sequential Evolving Algebras	7
3.1	Basic Transition Rules	7
3.1.1	Update Instructions	7
3.1.2	Two Rule Constructors	8
3.1.3	Guarded Multi-updates	8
3.2	Importing New Elements	9
3.2.1	Reserve	9
3.2.2	Transition Rules: Syntax	10
3.2.3	Auxiliary Vocabularies	11
3.2.4	Transition Rules: Semantics	11
3.2.5	Importing Several Elements at a Time	12
3.3	Programs and Runs	13
3.3.1	Programs and Pure Runs	13
3.3.2	External Functions	13

*©1995 Oxford University Press. Published in *Specification and Validation Methods*, Ed. E. Börger, Oxford University Press, 1995, 231–243.

[†]Partially supported by ONR grant N00014-91-J-1861 and NSF grant CCR-92-04742.

[‡]10 years later:

In the meantime, evolving algebras were renamed abstract state machines (ASMs).

For the reader's convenience, we added the table of contents and this footnote.

Yuri Gurevich, July 2005.

4	Nondeterministic Sequential Ealgebras and Some Other Simple Extensions of the Basic Model	14
4.1	Basic Evolving Algebras with Choice	14
4.1.1	Syntax	14
4.1.2	Semantics	15
4.1.3	Abbreviations	16
4.2	Some Other Simple Extensions of the Basic Model	16
4.2.1	First-order Guards	16
4.2.2	Qualified Choose Construct	17
4.2.3	Duplication	18
5	Parallelism: Evolving Algebras with Variables	18
5.1	Variables	18
5.2	Terms and Guards	19
5.3	A Parallel Version of the Basic EA Model	19
5.3.1	Syntax	19
5.3.2	Semantics of Rules	19
5.4	Importing Elements	20
5.5	Runs	21
6	Distributed Evolving Algebras	21
6.1	The Self Function	21
6.2	Basic Definition of Distributed Ealgebras	22
6.3	Generalizations	23
6.3.1	Active Agents	23
6.3.2	Active Teams	23
6.4	Sequential Runs	24
6.5	Partially Ordered Runs	24
6.5.1	Runs	24
6.6	Real-time Computations	25

1 Introduction

Computation models and specification methods seem to be worlds apart. The evolving algebra project started as an attempt to bridge the gap by improving on Turing’s thesis [5, 6]. We sought more versatile machines which would be able to simulate arbitrary algorithms in a direct and essentially coding-free way. Here the term algorithm is taken in a broad sense including programming languages, architectures, distributed and real-time protocols, *etc.*. The simulator is not supposed to implement the algorithm on a lower abstraction level; the simulation should be performed on the natural abstraction level of the algorithm.

The evolving algebra thesis asserts that evolving algebras are such versatile machines. The thesis suggests an approach to the notorious correctness problem that arises in mathematical modeling of non-mathematical reality: How can one establish that a model is faithful to reality? The approach is to construct an evolving algebra \mathcal{A} that reflects the given computer system so closely that the correctness can be established by observation and experimentation. (There are tools for running evolving algebras.) \mathcal{A} can then be refined or coarsened and used for numerous purposes. An instructive example is described in [1] by Egon Börger who championed this approach and termed \mathcal{A} the *ground model* of the system. The use of the successive refinement method is facilitated by the ability of evolving algebras to reflect arbitrary abstraction levels. This has been convincingly demonstrated by Börger and Rosenzweig in [4]; a simpler example is found in [7].

Evolving algebras have been used to specify languages (*e.g.* C, Prolog and VHDL), to specify real and virtual architectures (*e.g.* APE, PVM and Transputer), to validate standard language implementations (*e.g.* of Prolog, Occam), to validate distributed protocols (see examples in Parts III and IV of this book), to prove complexity results [2], *etc.*. See Börger’s annotated bibliography on evolving algebras in this book and the proceedings of the first evolving algebra workshop in [15].

Here we extend the definition of evolving algebras given in the tutorial [6] (henceforth “the tutorial”). For the sake of brevity, the term “evolving algebra” is often shortened to “ealgebra” (pronounced e-algebra) or “EA”; the latter term is used mostly as an adjective. Static algebras are discussed in §2. Sequential ealgebras are discussed in §3; first we define basic ealgebras and then we equip them with the ability to import new elements. Nondeterministic sequential ealgebras and some other simple extensions of basic ealgebras are discussed in §4, parallel ealgebras are discussed in §5, and distributed ealgebras are discussed in §6 which can be read immediately after §3. Admittedly this guide is harder to read than the tutorial, and we intend to write a more popular version of the guide.

Now let us return to the EA thesis. In the tutorial, we defined sequential ealgebras and sketched a speculative philosophical “proof” of the sequential version of the thesis. The definition of sequential ealgebras and the sequential EA thesis have survived several years of intensive application and experimentation. As a matter of fact, we (the EA community) seem to have run out of challenges.

The situation with non-sequential computations is more complicated. It seems that, for every reasonably understood class of algorithms, there is a natural extension of the basic EA model that “captures” that class. That form of the EA thesis also has survived several years of intensive application and experimentation. The philosophy and guiding principles of the EA approach seem quite stable. However, at the current stage of computer science, there is yet no clear understanding of what parallel, distributed or real-time algorithms are in general. Thus, the definitions of parallel and distributed ealgebras given below are necessarily tentative. They provide a foundation for existing EA applications and reflect my anticipation of things to come. (Many existing applications, including those in this volume, were done before this guide have been completed; the terminology there may reflect earlier versions of the guide.)

We try to derive our definitions from first principles. Unfortunately some arbitrariness is inescapable and one has to balance the clarity and simplicity versus programming convenience and efficient execution. When one thinks mostly about applications, as we do, there is a tendency to prefer programming convenience and efficient execution. This is a dangerous trend which leads to an idiosyncratic programming language. For future reference we formulate the following principle:

The Pragmatic Occam’s Razor Logic simplicity comes first; it may be sacrificed only in those cases where a slight logic complication is demonstrated to ease programming or improve execution efficiency in a substantial way.

The EA field is quickly expanding in depth and breadth. I hope that this guide lives up to its name and guides the developments in the near future.

Acknowledgment Egon Börger and Dean Rosenzweig generously shared with me their ideas and rich application experience. Discussions with Andreas Blass were indispensable in clarifying things. Numerous working walks with Jim Huggins through the woods of Ann Arbor were very helpful. Raghu Mani raised important implementation issues. Numerous ealgebraists commented on earlier drafts of the guide. I am very thankful to all of them. To an extent, this chapter is a result of a collective effort, though I am responsible for possible blunders.

A preliminary version of the guide has been tried out during the 1993 summer school on Specification and Validation Methods for Programming Languages and Systems on the beautiful island of Lipari in Italy. I use this opportunity to thank the organizers, Egon Börger and Alfredo Ferro, and all participants.

2 Static Algebras and Updates

2.1 Static Algebras: Motivation

In first-order logic, a structure is a nonempty set with operations and relations (called the basic operations and relations of the structure). That is how Tarski defined structures. He could have defined structures differently; there were a number of reasonable options. For our purposes here, a variant of Tarski’s notion is more appropriate. Respecting tradition, we do not redefine structures. Rather, we modify the notion of structure and give the new notion a new name.

Structures without relations are called *algebras* in the branch of mathematics called universal algebra. Restrict attention to algebras with distinct nullary operations *true* and *false* and define basic relations as basic operations taking only the Boolean values *true* and *false*. Further restrict attention to algebras with the equality relation and the usual Boolean operations. (We will specify later the values of the Boolean operations outside their natural domains.) The resulting notion of algebra is our variant of the notion of structure with equality. It allows us to write quantifier-free formulas as terms.

Actually, we are interested in multi-sorted structures with partial operations. The sorts can be given by unary relations (they will be called universes and the whole underlying set of a structure will be called the superuniverse). To deal with partial functions, further restrict attention to algebras with a nullary operation *undef*, different from *true* and *false*, and interpret an operation *f* as undefined at a tuple \bar{a} if $f(\bar{a}) = \text{undef}$. These algebras will be called *static algebras* or *states*. Their operations will be called functions.

In the following subsections, we start anew and define static algebras from scratch, establishing terminology on the way.

2.2 Vocabularies

A *vocabulary* (or *signature*) is a finite collection of function names, each of a fixed arity. Some function names may be marked as *relation names* or *static names*, or both. Every vocabulary contains the following static names: the equality sign, nullary function names *true*, *false*, *undef* and the names of the usual Boolean operations. The equality sign and *true*, *false* are marked as relation names. The Greek letter Υ is reserved to denote vocabularies.

Logic Names The particular function names listed above are *basic logic names*. There are precedents of logic names in mathematical logic, though usually they are called logical constants. For example, the equality sign is a logic name in first-order logic with equality. Usually, logic names are present in every vocabulary and their interpretations satisfy some *a priori* restrictions. Accordingly, we suppose that the basic logic names appear in every vocabulary, and thus there is no need to mention them when a particular vocabulary is described.

An additional logic name is introduced in §3. It does not necessarily appear in every vocabulary and it is not marked static. The latter is one reason why we do not use the term “logical constants”.

2.3 Definition of Static Algebras

A *static algebra* or (for the sake of brevity) *state* S of vocabulary Υ is a nonempty set X , the *superuniverse* of S , together with interpretations of the function names in Υ on X . An r -ary function name is interpreted as a function from X^r to X , a *basic function* of S . The interpretation of an r -ary relation name is a function from X^r to $\{true, false\}$, a *basic relation* of S . The vocabulary Υ is called the *vocabulary* of S and denoted $\text{Fun}(S)$.

The interpretations of the nullary logic names *true*, *false* and *undef* are distinct elements of X . The Boolean operations behave in the usual way on the Boolean values *true* and *false* and produce *undef* if at least one of the arguments is not Boolean. The equality sign is interpreted as the characteristic function of the identity relation on X . If $f(\bar{x})$ evaluates to *true* in S , we say that $f(\bar{x})$ holds in S ; and if $f(\bar{x})$ evaluates to *false* in S , we say that $f(\bar{x})$ fails in S .

Formally speaking, basic functions are total. However, we view them as being partial and define the domain $\text{Dom}(f)$ of an r -ary basic function f as the set of r -tuples \bar{x} such that $f(\bar{x}) \neq \text{undef}$. Let us stress though that *undef* is an ordinary element of the superuniverse. Often, a basic function produces *undef* if at least one argument equals *undef*, but this is not required and there are exceptions (*e.g.* basic relations).

Universes A basic relation f may be viewed as the set of tuples where it evaluates to *true*. We may write $\bar{x} \in f$ instead of $f(\bar{x})$. If f is unary it can be viewed as a special *universe*. For example, we may have a universe *Nodes* and declare a binary relation *Edge* over the universe of *Nodes*; $\text{Edge}(x, y)$ will hold only if both x and y belong to *Nodes*. Such universes allow us to view states as many-sorted structures. Sometimes we speak about universe names. These are unary relation names intended to be used as universes.

As a rule, *undef* is not included in universes. Coming back to our example, is it natural that $\text{Edge}(\text{undef}, \text{undef})$ equals *false* rather than *undef*? In a sense, yes. Think about *Edge* as a set of pairs of nodes. It is natural that the pair $(\text{undef}, \text{undef})$ does not belong there.

2.4 Terms

Terms are defined recursively, as in first-order logic:

- A variable is a term.
- If f is an r -ary function name and t_1, \dots, t_r are terms, then $f(t_1, \dots, t_r)$ is a term.

As usual, *ground terms* are terms without variables. By analogy, other syntactical objects without variables will be called ground.

Atomic Boolean terms are terms of the form $f(\bar{t})$, where f is a relation name. *Boolean terms* are built from atomic Boolean terms by means of the Boolean operations.

Appropriate States and the Fun Notation In addition to terms, we will define various other syntactic objects, *e.g.*, update instructions and transition rules. We call a state S *appropriate* for a syntactic object s if $\text{Fun}(S)$ includes the collection of function names that occur in s . By default (that is, unless explicitly defined differently), that collection will be denoted $\text{Fun}(s)$.

In an appropriate state S , a ground term $t = f(t_1, \dots, t_r)$ evaluates to an element $\text{Val}_S(t) = f(\text{Val}_S(t_1), \dots, \text{Val}_S(t_r))$. If \bar{t} is a tuple (t_1, \dots, t_r) of terms, define $\text{Val}_S(\bar{t}) = (\text{Val}_S(t_1), \dots, \text{Val}_S(t_r))$.

An expression $t_1 = t_2$ may be a Boolean term or a metalanguage statement. Often it does not matter which it is. One can use two different equality signs or just try to be careful; we choose the second alternative.

2.5 Locations and Updates

As in first-order logic, the *reduct* of an Υ -state S to a smaller vocabulary Υ' is the Υ' -state S' obtained from S by “disinterpreting” function names in $\Upsilon - \Upsilon'$; S is an *expansion* of S' to Υ .

A *carrier* is a state whose vocabulary contains only static function names. The *carrier* $|S|$ of a state S is the reduct of S to the static part of $\text{Fun}(S)$.

A *location over a carrier* C is a pair $\ell = (f, \bar{x})$, where f is a function name outside of $\text{Fun}(C)$ and \bar{x} is a tuple of elements of C whose length equals the arity of f ; location ℓ is *relational* if f is a relation symbol. $\text{Loc}_\Upsilon(C)$ is the collection of all locations over C with function names in Υ . An Υ -state S with carrier C will sometimes be viewed as a function from $\text{Loc}_\Upsilon(C)$ to (the superuniverse of) C ; *locations of* S are locations in $\text{Loc}_\Upsilon(C)$.

If a state S is appropriate for a ground term $t_0 = f(\bar{t})$, then the *location* of t_0 in S is the location $(f, \text{Val}_S(\bar{t}))$.

An *update* of a state S is a pair $\alpha = (\ell, y)$, where ℓ is a location of S and $y \in |S|$; if ℓ is relational then y is Boolean. (More precisely, y belongs to the superuniverse of static algebra $|S|$; the looser language is common in logic.) The location ℓ is the *location* $\text{Loc}(\alpha)$ of α , and y is the *value* $\text{Val}(\alpha)$ of α . To *fire* α at S , put y into the location ℓ ; that is, redefine S to map ℓ to y . The result is a new state S' such that $\text{Fun}(S') = \text{Fun}(S)$, $|S'| = |S|$, $S'(\ell) = y$ and $S'(\ell') = S(\ell')$ for every location ℓ' of S different from ℓ .

2.6 Update Sets and Families of Update Sets

An *update set* β over a state S is a set of updates of S . $\text{Loc}(\beta) = \{\text{Loc}(\alpha) : \alpha \in \beta\}$. For each $\ell \in \text{Loc}(\beta)$, $\text{Val}_\beta(\ell) = \{\text{Val}(\alpha) : \alpha \in \beta \wedge \text{Loc}(\alpha) = \ell\}$.

An update set β is *consistent* at the given state S if every $\text{Val}_\beta(\ell)$ is a singleton set; otherwise β is inconsistent.

To fire a consistent β at the given state S , fire all its members simultaneously. The result is a new state S' with the same vocabulary and carrier as S . If $\ell \in \text{Loc}(\beta)$ then $S'(\ell)$ is the only

element of $\text{Val}_\beta(\ell)$; otherwise $S'(\ell) = S(\ell)$. To fire an inconsistent update set β at the given state S , do nothing; the new state S' equals S .

Remark It is reasonable to require that the detection of inconsistency manifest itself in some way; for example, a nullary function *crash* automatically gets value *true*. To keep the EA logic clean and simple, we try to minimize the number of things done automatically, and thus we leave necessary manifestations of inconsistency to the programmer. This is one application of the pragmatic Occam's razor of §1; substantial programming convenience has not been demonstrated yet.

To fire a family γ of update sets over S , nondeterministically choose some update set $\beta \in \gamma$ and fire it at S . If $\gamma = \emptyset$, do nothing. Intentionally, the empty family of update sets means inconsistency.

2.7 Conservative Determinism vs. Local Nondeterminism

The mode of dealing with inconsistent update sets described above can be called conservative determinism. The mode of dealing with inconsistent update sets in the tutorial was different: Fire all updates simultaneously; in case of conflict at any location ℓ , choose the new value for ℓ nondeterministically among all candidate values. It could be called local nondeterminism.

With the exception of this change in the treatment of inconsistent update sets, this guide is compatible with the tutorial. The change is not as big as it may seem because people are usually interested in deterministic programs. As far as we know, no existing EA application is affected. The local nondeterminism has not been exploited. The conservative determinism is simpler, and a more manageable form of nondeterminism will be introduced in §4.

3 Sequential Evolving Algebras

Basic transition rules are defined in subsection 3.1. Subsection 3.2 deals with the problem of extending universes. The reader may skip 3.2 and go directly to subsection 3.3 on programs and runs.

3.1 Basic Transition Rules

In this subsection, terms are ground.

3.1.1 Update Instructions

An *update instruction* R is an expression

$$f(\bar{t}) := t_0$$

where f is a non-static function name (the *subject* of the instruction), \bar{t} is a tuple of terms whose length equals the arity of f , and t_0 is another term; if f is a relation name then t_0 must be a Boolean term. (Update instructions are called local function updates in the tutorial.)

Semantics To execute R at an appropriate state S , fire the update $\alpha = (\ell, y)$ at S , where $\ell = (f, \text{Val}_S(\bar{t}))$ and $y = \text{Val}_S(t_0)$. For future reference define $\text{Updates}(R, S) = \{\alpha\}$.

3.1.2 Two Rule Constructors

Basic rules are constructed recursively from update instructions by means of two rule constructors: the sequence constructor and the conditional constructor. Semantics is defined by means of update sets. For each rule R and every state S appropriate for R , we define an update set $\text{Updates}(R, S)$ over S . To fire R at S , fire $\text{Updates}(R, S)$.

The Sequence Constructor A sequence of rules is a rule.

Semantics If R is a sequence of rules R_1, \dots, R_k then

$$\text{Updates}(R, S) = \text{Updates}(R_1, S) \cup \dots \cup \text{Updates}(R_k, S).$$

In other words, to fire a sequence of rules, fire all of them simultaneously. Notice that $\text{Updates}(R, S)$ is inconsistent if any R_i is so.

Remark The term “sequence” may be misleading here. We are not executing first R_1 , then R_2 , then R_3 , etc.. A better term is “block”. (This remark is written at the proofreading stage.)

The Conditional Constructor If k is a natural number, g_0, \dots, g_k are Boolean terms and R_0, \dots, R_k are rules, then the following expression is a rule:

```

if  $g_0$  then  $R_0$ 
elseif  $g_1$  then  $R_1$ 
  :
elseif  $g_k$  then  $R_k$ 
endif

```

If the guard g_k is the nullary function *true*, then the last elseif clause may be replaced by “else R_k ”. For brevity we will say that the conditional rule R above is the conditional rule with clauses $(g_0, R_0), \dots, (g_k, R_k)$.

Semantics $\text{Updates}(R, S) = \text{Updates}(R_i, S)$ if g_i holds in S but every g_j with $j < i$ fails in S . $\text{Updates}(R, S) = \emptyset$ if every g_i fails in S .

3.1.3 Guarded Multi-updates

A *multi-update instruction* is a sequence of update instructions. A *guarded update instruction* (respectively, *guarded multi-update instruction*) is a rule of the form

```

if  $g$  then  $R$  endif

```

where R is an update (respectively, a multi-update) instruction.

Lemma 3.1 For every rule R , there is a sequence R' of guarded updates such that $\text{Fun}(R') = \text{Fun}(R)$ and $\text{Updates}(R', S) = \text{Updates}(R, S)$ for all appropriate states S .

For example, the rule

```
if FirstChild(c) ≠ undef then c:=FirstChild(c)
elseif NextSib(c) ≠ undef then c:=NextSib(c)
elseif Parent(c) ≠ undef then c:=Parent(c)
endif
```

converts to the following sequence of guarded updates:

```
if FirstChild(c) ≠ undef then c:=FirstChild(c) endif
if FirstChild(c)=undef and NextSib(c) ≠ undef then
  c:=NextSib(c) endif
if FirstChild(c)=undef and NextSib(c)=undef
  and Parent(c) ≠ undef then c:=Parent(c) endif
```

The Lemma suggests a simpler definition of rules. The reason for choosing the recursive definition is pragmatic. It is too tedious to write rules as sequences of guarded updates. It is feasible to write them as sequences of guarded multi-updates but it is more convenient and practical to use `elseif` clauses and nest conditionals. The pragmatic Occam's razor does not cut as much as the original Occam's razor would.

Remark This is another proofreading time remark. The new version of the EA interpreter permits the use of two additional rule constructors. One is the case constructor, like that in Pascal, which may make the execution substantially more efficient. Of course, the same set of updates is generated by a case command and its case-free equivalent; the difference is in how fast this set is generated. For example, consider a sequence of rules of the form “if $t = i$ then R_i endif” where i ranges from 1 to a relatively large n . This example is extreme, because the the set $\{1, \dots, n\}$ of alternatives is so easy to deal with; but it is not unusual to have similar long sequences of rules. In addition, the case construct makes it easier to program a sequential execution of a sequence of rules, which is sometimes desirable. The other rule constructor is “let $x=t$ in R ”, which prevents re-evaluations of term t in R and which has been used informally. The let constructor was advocated by Raghu Mani who is working on the new EA interpreter.

3.2 Importing New Elements

The basic rules suffice for many purposes (*e.g.*, for describing the C programming language [7]), but they do not suffice to model all sequential algorithms. A sequential algorithm may add a new node to a graph or create a new message. We need rules that allow us to create new nodes, new messages, *etc.*, and such rules are introduced in this subsection. However, we do not create new elements; instead, we use a special universe Reserve from which the new elements come.

In this section we use individual variables, but only in a limited way. (Variables are used more extensively in §5.) Roughly speaking, only bound variables are used; free variables appear only in contexts where some values have been assigned to them.

3.2.1 Reserve

In addition to basic logic names, we introduce a new logic name: a universe name Reserve. It is not static, and we do not require that it belong to the vocabulary of every static algebra. If the

vocabulary of state S contains `Reserve`, then the set $\{x : S \models x \in \text{Reserve}\}$ is the *reserve* of S . Intuitively the reserve is a naked set.

Reserve Proviso Every state satisfies the following conditions:

- Every basic relation, with the exception of equality and `Reserve`, evaluates to *false* if at least one of its arguments belongs to the reserve.
- Every other basic function evaluates to *undef* if at least one of its arguments belongs to the reserve.
- No basic function outputs an element of the reserve.

It follows that every permutation of the reserve is an automorphism of the state.

3.2.2 Transition Rules: Syntax

Generalize the definitions of terms and update instructions in 3.1 as follows:

- allow terms to have variables, and
- forbid mentioning `Reserve`.

Variables are often treated as auxiliary nullary function names below but *a variable cannot be the subject of an update instruction*. The reason for forbidding to mention `Reserve` in terms and update instructions is discussed below.

Rules are constructed from update instructions by means of three rule constructors: the sequence constructor, the conditional constructor and the import constructor.

The Import Constructor If v is a variable and R_0 is a rule, then the following expression is a rule with *main existential variable* v and *body* R_0 :

```
import v
  R0
endimport
```

In the usual and obvious way define which occurrences of variables are free and which are bound. Call a rule *perspicuous* if no variable has both bound and free occurrences, and no bound variable is declared more than once. (The latter means here that different occurrences of the import command have different main existential variables.)

Let $\text{Free}(R)$ be the set of free variables of a rule R . In other words, $\text{Free}(R)$ is the set of variables v such that v occurs freely in rule R . Define $\text{Bound}(R)$ similarly. If R is an import rule with main existential variable v and body R_0 , we have:

$$\text{Free}(R) = \text{Free}(R_0) - \{v\}, \quad \text{and} \quad \text{Bound}(R) = \text{Bound}(R_0) \cup \{v\}.$$

3.2.3 Auxiliary Vocabularies

The names of variables are different from function names of course, but it is convenient to treat free variables of rules as auxiliary nullary functions (which cannot be subjects of update instructions). An *auxiliary vocabulary* has the form $\Upsilon \cup V$, where Υ is a genuine vocabulary and V is a finite set of variables.

If S is a state of an auxiliary vocabulary $\Upsilon' = \Upsilon \cup V$, then $\text{Fun}(S) = \Upsilon'$. S is *appropriate* for a rule R if Υ contains all function names of R and V contains all free variables of R . R is *S-perspicuous* if it is perspicuous and its bound variables do not occur in V .

3.2.4 Transition Rules: Semantics

An import commands chooses an element of the reserve and removes it from the reserve. To clarify our intentions, we note that the non-perspicuous rule

```
import v
  Parent(v) := CurrentNode
endimport
import v
  Parent(v) := CurrentNode
endimport
```

creates *two* children of CurrentNode. In general, different choices from the reserve produce different elements.

For each rule R and every state S appropriate for R , we define an update set $\text{Updates}(R, S)$ over S ; to fire R at S , fire $\text{Updates}(R, S)$.

First, we consider the case of when R is S -perspicuous. Fix an injective map ξ from $\text{Bound}(R)$ to the reserve of the given S . (The injectivity means that ξ assigns different elements to different bound variables.) By induction on subrule R' of R we define sets $\text{Updates}(R', S', \xi)$ where S' is an expansion of S appropriate for R' and such that R' is S' -perspicuous. (Recall that S' is an expansion of S if and only if the reduct of S' to $\text{Fun}(S)$ equals S .) Let $\Upsilon' = \text{Fun}(S')$.

The cases of update instructions, sequence rules and conditional rules are treated as above. (Variables in Υ' are treated as nullary functions.) Suppose that R' is an import rule with main existential variable v and body R_0 . Let $a = \xi(v)$ and S'_a be the expansion of S' to the auxiliary vocabulary $\Upsilon' \cup \{v\}$ where v is interpreted as a . Recall that variables are not subjects of update instructions. Thus $\text{Updates}(R_0, S'_a, \xi)$ is an update set over S' . Set

$$\text{Updates}(R', S', \xi) = \{((\text{Reserve}, a), \text{false})\} \cup \text{Updates}(R_0, S'_a, \xi).$$

Finally $\text{Updates}(R, S) = \text{Updates}(R, S, \xi)$. Of course, $\text{Updates}(R, S)$ is not defined uniquely, because it depends on ξ . It is easy to see, however, the resulting state is unique up to isomorphism.

Second, we stipulate that an arbitrary rule R is equivalent, over the given appropriate state S , to an S -perspicuous rule R' obtained from R by renaming the bound variables. (The desired R' can be obtained by iterating the following transformation: Select an innermost import subrule R_1 whose main existential variable v occurs in the rest of the rule or in $\text{Fun}(S)$, and replace v with a fresh variable in R_1 .) The stipulation means the following: To fire R at S , fire R' at S .

Discarding Elements from Universes Finally, we explain the reason for forbidding to mention Reserve explicitly in our rules. Terms $\text{Reserve}(t)$ always evaluate to *false*, so evaluating $\text{Reserve}(t)$

or setting it to *false* is useless. But why not to allow putting the value *true* into Reserve locations. Elements can be discarded from universes, of course; to discard an element (represented by a term) t from a universe U , use the instruction $U(t) := \text{false}$. Isn't the reserve a natural place for unwanted elements? Yes, it is. Notice, however, that moving an element into the reserve may necessitate numerous changes of basic functions in order to ensure that the Reserve proviso remains valid. Would such a move contradict the sequential character of our rules? Not necessarily. We could just mark discarded elements as reserve elements, but then it might be necessary to augment rules with numerous guards $\text{Reserve}(t) = \text{false}$, which would be too tedious. It is preferable to leave the discarded elements alone. This pragmatic argument was put forward originally by Egon Börger.

But shouldn't the computational resources of the ealgebra simulating an algorithm A closely reflect the computational resources of A ? Yes, but it is important to separate the following concerns: the logic of A and the relevant resources of A . Concentrating on the logic of A may allow one to come up with simpler rules for the simulating ealgebra. And if one needs to track the resources of A , a separate bookkeeping may be set up. This separation of concerns allows us, for example, to use infinite universes. And caring about only particular elements and universes, rather than the whole superuniverse, makes combining ealgebras easier.

3.2.5 Importing Several Elements at a Time

Let v_1, v_2 be distinct variables. Abbreviate

<pre>import v1 import v2 R0 endimport endimport</pre>	to	<pre>import v1, v2 R0 endimport</pre>
---	----	---

In a similar way, define abbreviations

```
import v1, ..., vk
  R0
endimport
```

Abbreviate

<pre>import v1, ..., vk U(v1) := true ⋮ U(vk) := true R0 endimport</pre>	to	<pre>extend U with v1, ..., vk R0 endextend</pre>
--	----	---

Later (in 5.4) we'll see how to import a number of elements that is not bounded *a priori* by any constant. Here is an example of the extend rule:

```

extend Nodes with  $v_1, v_2$ 
  FistChild(CurrentNode) :=  $v_1$ 
  SecondChild(CurrentNode) :=  $v_2$ 
  NextSib( $v_1$ ) :=  $v_2$ 
endextend

```

3.3 Programs and Runs

3.3.1 Programs and Pure Runs

A *program* P is a rule without free variables. A *basic program* is a basic rule without free variables. In applications, a program is usually a sequence of rules referred to as rules of the program. To fire P at an appropriate state S , fire $\text{Updates}(P, S)$ at S .

A *pure run* of P is a sequence $\langle S_n : n < \kappa \rangle$ of states of vocabulary $\text{Fun}(P)$ such that each S_{n+1} is obtained from S_n by firing P at S_n . Here and henceforth κ is a positive integer or the first infinite ordinal. In the latter case, $\{n : n < \kappa\}$ is the set of all natural numbers.

The adjective “pure” reflects the fact that the run is not affected by the environment.

3.3.2 External Functions

In general runs may be affected by the environment. Suppose that the environment manifests itself via some basic functions e_1, \dots, e_k , called *external functions*. A typical external function is the input provided by the user.

Think about an external function as a (dynamic) oracle. The ealgebra provides the arguments and the oracle gives the result. The oracle need not be consistent and may give different results for the same argument at different times. The seeming inconsistency may be quite natural. For example, the argument may specify an input channel. The next time around, another input can come via the same channel.

However, the oracle should be consistent during the execution of any one step of the program. In an implementation, this may be achieved by not reiterating the same question during a one-step execution. Ask the question once and, if necessary, save the result and reuse it.

The computation steps of a program are supposed to be atomic at an appropriate level of abstraction. A computation step is hardly atomic if during that step the ealgebra queries an oracle and then, depending on the result, submits another query to the same or a different oracle. Thus it seems reasonable to forbid nesting of external functions. Indeed, the need to nest external functions has not arisen in applications so far. But we withhold final judgement and wait for more experimentation.

Call non-external basic functions *internal*. If S is an appropriate state for a program P , let S^- be the reduct of S to the internal vocabulary.

Runs A *run* of a program P is a sequence $\langle S_n : n < \kappa \rangle$ of states where:

- every nonfinal S_n is an appropriate state for P and the final state (if any) is a state of the internal vocabulary of P , and
- every S_{n+1}^- is obtained from S_n by firing P at S_n .

Internal and External Locations It may happen that the environment controls only a part of a function e_i and the remaining part of e_i is governed internally. In such a case it is natural to speak about internal and external locations rather than internal and external functions. See an example in [3, 3.1]. The generalization to that case is relatively straightforward.

Irrelevant Values of External Functions In order to fire a given program at a given state, we may not need to know all about the state. Only some values of external functions may be needed for firing. We may not care about or even know the values of external functions which are not needed for the execution. Some of those values may even be ill-defined. There is also an issue of influencing the environment by requiring an extra value, *e.g.*, by requiring a user-provided datum.

It is natural to set all irrelevant values of external functions to *undef*. However, caution should be exercised in the distributed situation (see §6) where other agents may have different views of those values.

Sometimes it may be simpler to use partial states. A partial Υ -state S with carrier C can be defined as a partial function from $\text{Loc}\Upsilon(C)$ to C . See examples in [3, 8]. For simplicity, we will not use partial states here.

4 Nondeterministic Sequential Ealgebras and Some Other Simple Extensions of the Basic Model

Describing algorithms on higher abstraction levels, one often comes across the phenomenon of nondeterminism. Nevertheless, the built-in nondeterminism of ealgebras has been rarely used. It is often more appropriate to use external functions to reflect nondeterministic behavior. (In the distributed case, nondeterminism may be often eliminated by introducing additional agents.) Consider for instance the assignment statement of the C programming language. Should one evaluate the left side or the right side first? According to the ANSI standard (ANSI is the American National Standards Institute), the choice of the evaluation order is implementation-dependent. Moreover, an implementation does not have to be consistent; the evaluation order may change when the same assignment statement is executed next time around (say, in a loop). This is an obvious case of nondeterminism and first we, the authors of [7], were tempted to use a nondeterministic rule to reflect the nondeterminism. But then we realized that C is perfectly deterministic. It is just that execution may depend on information provided by implementation. Thus it is more faithful to the standard (and more convenient) to use an external function that decides the evaluation order.

Still, nondeterministic commands may be desired and we provide such commands in this section. For example, it may be convenient to formalize the environment in a distributed situation, so that an external function of one agent is nondeterministically computed by another agent.

For simplicity, we ignore the import constructor in this section. It is easy to extend the language of this section with the import constructor. Moreover, the choice constructor defined below and the import constructor can be combined into one constructor.

4.1 Basic Evolving Algebras with Choice

4.1.1 Syntax

Transition rules are constructed as in 3.2, except that instead of the import constructor, we use the Choose (or Choice) Constructor:

Choose Constructor If U is a universe name different from Reserve, v is a variable and R_0 is a rule then the following expression is a rule with *main existential variable* v that ranges over U and *body* R_0 :

```
choose  $v$  in  $U$ 
   $R_0$ 
endchoose
```

This is the basic version of the choice constructor; a stronger version is defined in 4.2.2. Perspicuity is defined as 3.2.

4.1.2 Semantics

For each rule R and each state S appropriate for R , we define a family $\gamma = \text{NUpdates}(R, S)$ of update sets over S . To fire R at S , choose any $\beta \in \gamma$ and fire β at S .

We stipulate that an arbitrary rule R is equivalent, over the given S , to an S -perspicious rule R' obtained from R by renaming the bound variables. The equivalence means here that $\text{NUpdates}(R, S) = \text{NUpdates}(R', S)$. It remains to define $\gamma = \text{NUpdates}(R, S)$ when S is S -perspicious.

Global Choice Semantics Semantics is defined as in 3.2.4. On one hand, things are simpler this time around because there is no correlation among individual choices. On the other hand, there is a complication related to attempts to choose an element of the empty set. Such attempt cannot succeed and the execution should be aborted. To deal with this complication, we extend the collection of updates of any state by an ideal element \perp that symbolizes inconsistency. If an update set β contains \perp then firing β does not change the state; we call such β contradictory.

Suppose that a state S is appropriate for a rule R and R is S -perspicious. Let V be the collection of bound variables of R such that the range of v is not empty in state S . Fix a function ξ on V such that, for each $v \in V$, $\xi(v)$ belongs to the range of v in S . By induction on subrule R' of R define $\text{Updates}(R', S', \xi)$ where S' is an expansion of S appropriate for R' and R' is S' -perspicious.

The cases of update instructions, sequence rules and conditional rule are treated as above. Notice that if R' is a sequence of rules R_i and some $\text{Updates}(R_i, S', \xi)$ is contradictory then $\text{Updates}(R', S', \xi)$ is so.

Suppose that R' is a choose rule with main existential variable v and body R_0 . If the range of v is empty then $\text{Updates}(R', S', \xi) = \perp$. Otherwise let $a = \xi(v)$ and S'_a be the expansion of S' to the auxiliary vocabulary $\Upsilon' \cup \{v\}$ where v is interpreted as a . Set $\text{Updates}(R', S', \xi) = \text{Updates}(R_0, S'_a, \xi)$.

Finally, $\text{NUpdates}(R, S)$ is the set of $\text{Updates}(R, S, \xi)$ where ξ takes all possible values.

Semantics without Global Choice The global choice semantics is straightforward. However, contrary to the situation 3.2.4, there is no correlation among individual choices this time around, and thus there is no real need for a global choice function ξ . It may be more elegant to define $\gamma = \text{NUpdates}(R, S)$ directly by induction on R . We suppose again that S is appropriate to R and R is S -conspicuous.

If R is an update instruction then $\gamma = \{\text{Updates}(R, S)\}$. If R is a sequence of rules R_1, \dots, R_k , then

$$\gamma = \{\beta_1 \cup \dots \cup \beta_k : \text{each } \beta_i \in \text{NUpdates}(R_i, S)\}.$$

Notice that γ is empty if some so is $\text{NUpdates}(R_i, S)$.

If R is a conditional rule with clauses $(g_0, R_0), \dots, (g_k, R_k)$, we have two cases as usual; if all $k + 1$ guards fail in S then $\gamma = \{\emptyset\}$, and if g_i is the first guard that holds in S then $\gamma = \text{NUpdates}(R_i, S)$. (It would be a mistake to replace $\{\emptyset\}$ with \emptyset above. If $\text{NUpdates}(R_1, S) = \emptyset$ then $\text{NUpdates}((R_1, R_2), S) = \emptyset$ for every rule R_2 , which is not desired.)

Finally, suppose that R is a choose rule with universe name U , main existential variable v and body R_0 . For each $a \in U$, let S_a be the expansion of S of the auxiliary vocabulary $\text{Fun}(S) \cup \{v\}$ where v is interpreted as a . Then

$$\gamma = \bigcup \{\text{NUpdates}(R_0, S_a) : a \in U\}.$$

Notice that γ is empty if U is empty.

It is easy to check that if R contains no choice subrules then $\text{NUpdates}(R, S) = \{\text{Updates}(R, S)\}$.

Remark In the second approach, \perp is not used. Its role is played by the empty family of update sets. This gives us an idea to eliminate the use of \perp in the first approach: replace $\text{Updates}(R', S', \xi)$ with the singleton family $\{\text{Updates}(R', S', \xi)\}$ and replace \perp with the empty family.

Runs The definition of runs in §3 remains in force.

4.1.3 Abbreviations

Let v_1, v_2 be distinct variables. Abbreviate

<pre> choose v_1 in U choose v_2 in U R_0 endchoose endchoose </pre>	to	<pre> choose v_1, v_2 in U R_0 endchoose </pre>
---	----	--

In a similar way define abbreviation

```

choose  $v_1, \dots, v_k$  in  $U$ 
   $R_0$ 
endchoose

```

4.2 Some Other Simple Extensions of the Basic Model

We consider three extensions, which are simple in the sense that it is easy to define them. The third extension has not been used; it is just a trial balloon.

4.2.1 First-order Guards

In §3, guards were Boolean terms. Now we introduce a separate syntactic category of guards. Intuitively, guards are first-order formulas with bound variables. It is intended that bound variables range over finite domains, though exceptions are possible. Here is a recursive definition:

- If f is an r -ary relation name and t_1, \dots, t_r are terms, then $f(t_1, \dots, t_r)$ is a guard.

- Any Boolean combination of guards is a guard.
- If g is a guard and U a universe name, then $(\exists v \in U)g$ and $(\forall v \in U)g$ are guards.

Call a guard closed if it has no free variables. Extend the definition of basic ealgebras by replacing the condition “ g_1, \dots, g_k are Boolean terms” with the condition “ g_1, \dots, g_k are closed guards” in the definition of the conditional rule constructor.

Semantics The definition of the value of a closed guard at an appropriate state mirrors the truth definition of formulas in first-order logic. The semantics of rules is given exactly as in 3.1.

Remark One can go further in this direction and use quantification inside other terms. To formalize this idea, the notion of terms can be redefined as follows:

- A variable v is a term.
- If f is an r -ary function name and t_1, \dots, t_r are terms, then $f(t_1, \dots, t_r)$ is a term. The new term is Boolean if f is a relation name.
- Boolean terms are closed under the Boolean operations and quantification, and every Boolean term is a term.

4.2.2 Qualified Choose Construct

Restricting the choice by a Boolean term gives a much more powerful version of the choose constructor.

Qualified Choose Constructor If U is a universe name different from Reserve, v is a variable, $g(v)$ is a Boolean term and R_0 is a rule, then the following expression is a rule with *main existential variable* v that ranges over U and body R_0 :

```
choose  $v$  in  $U$  satisfying  $g(v)$ 
   $R_0$ 
endchoose
```

Replacing the choose constructor with the qualified choose constructor requires only a small and obvious change in the semantical definition of 4.1.2. We restrict attention to the global choice approach. Consider the case in the inductive definition of $\text{Updates}(R', S', \xi)$ where R' is a choose rule and the range U of the main existential variable v of R in S' is not empty. If $g(\xi(v))$ fails in S , set $\text{Updates}(R', S', \xi) = \perp$.

It is easy to construct a rule to choose several elements v_1, \dots, v_k subject to a condition $g(v_1, \dots, v_k)$.

The qualified choose constructor may be too powerful. The decision problem whether there is any tuple (v_1, \dots, v_k) in the universe U satisfying the condition g may be hard. If U is the set of natural numbers and g a polynomial, the decision problem may even be undecidable [13]. But the logical clarity of the constructor is attractive. It may be used in particular to reflect environmental forces that are not necessarily algorithmic.

4.2.3 Duplication

The powerful extension of basic ealgebras considered in this subsection is logically clear but untried and computationally expensive. It does not hurt to explore it though.

Call elements a and a' of a state S indistinguishable as arguments for a basic r -ary function f if $f(b_1, \dots, b_r) = f(c_1, \dots, c_r)$ for all r -tuples b_1, \dots, b_r and c_1, \dots, c_r such that either $b_i = c_i$ or $\{b_i, c_i\} = \{a, a'\}$. Call a, a' indistinguishable as arguments if they are indistinguishable as arguments for any basic function with the exception of equality. Now we are ready to introduce the duplicate constructor:

```
duplicate  $t$  as  $v$ 
   $R_0$ 
endduplicate
```

Semantics To execute, calculate $a = \text{Val}_S(t)$, get some a' from the reserve and redefine basic functions on tuples involving a' in such a way that a and a' become indistinguishable as arguments. Then execute R_0 with v equal a' .

Duplication can be seen as a powerful inheritance mechanism. It is easy to see that the extend construct is not powerful enough to replace duplication.

5 Parallelism: Evolving Algebras with Variables

What does it mean that an algorithm is sequential? This usually means that the algorithm has the following two features. First, time is sequential. The algorithm proceeds from some initial state S_0 to a state S_1 , then to a state S_2 , *etc.*, and the steps are atomic. Second, only a bounded amount of work is done at each step. In principle, a single agent is able to move the algorithm from S_0 to S_1 , then to S_2 , *etc.*

In this section, we are interested in one-agent algorithms where the agent may perform a substantial amount of work at one step. We use variables to formalize such algorithms. It is intended that non-Reserve variables range over finite (better yet, feasible) domains, though exceptions are possible.

We do not assume any particular sequential order of executing one step of the algorithm. It is possible that this work involves plenty of parallelism and is implemented by a number of auxiliary agents. But on the natural level of abstraction of the given algorithm, those auxiliary agents are invisible, and in principle a single agent may execute the algorithm.

5.1 Variables

In preceding sections, we dealt with implicit variables declarations by means import commands, bounded quantifiers, etc. In this section, we introduce explicit variable declarations.

An *explicit atomic variable declaration* is an expression “Var v ranges over U ”, where v is a variable and U a universe name. The universe U is the range (or type) of the variable v . A *explicit variable declaration* D is a sequence of explicit atomic variable declarations, and $\text{Var}(D)$ is the collection of variables in D . For brevity, the adjective explicit is often omitted.

Intuitively, D is a set of explicit atomic declarations, but we do not forbid re-declarations of the same variable. The range of a variable $v \in \text{Var}(D)$ is the range in the last declaration of v in

D . In other words, later declarations of a variable override the earlier ones. One may use more concise explicit variable declarations, like “Var v_1, \dots, v_k range over U ”.

A variable declaration D *covers* a syntactic object s if $\text{Var}(D)$ contains all free (that is undeclared) variables of s .

As in 3.2.3, we use auxiliary vocabularies of the form $\Upsilon \cup V$, where Υ is a genuine vocabulary, V a finite set of variables and each $v \in V$ is treated as a nullary function, except it cannot be the subject of an update instruction. We say that a state S of an auxiliary vocabulary is *appropriate* for a syntactical object s if all function names and all free variables of s occur in $\text{Fun}(S)$.

5.2 Terms and Guards

Terms and Boolean terms are defined in §3. Guards are defined in 4.2.1. The free variables of terms and guards are defined inductively, as in first-order logic. Notice that a bounded quantifier implicitly contains an atomic declaration.

As usual, every guard g is equivalent to a guard g' where no variable is both bound and free and where different quantifier occurrences bind different variables. To reduce g to g' , iterate the following transformation: Select an innermost quantifier q whose variable v occurs outside the scope of q and then replace v with a fresh variable in the scope of q .

5.3 A Parallel Version of the Basic EA Model

5.3.1 Syntax

Update instructions and basic rules are defined as in 3.1, except that terms may have free variables, and guards are defined as above. In addition, we have the following third rule constructor.

The Declaration Constructor An atomic variable declaration followed by a rule is a rule.

By an obvious induction on rules, define which occurrences of variables are free (or undeclared) and which are bound. Suppose that D is a variable declaration, R is a rule, and S is a state of an auxiliary vocabulary. R is (D, S) -*perspicuous* if it satisfies the following conditions:

- no variable is declared (explicitly or implicitly) more than once in R , and
- $\text{Bound}(R)$ is disjoint from $\text{Free}(R) \cup \text{Var}(D) \cup \text{Fun}(S)$.

Programs A *program* is a rule without any undeclared variables.

5.3.2 Semantics of Rules

By induction on R , we define the update set $\beta = \text{Updates}(D, R, S)$ generated by a rule R at an appropriate state S under a declaration D that covers R . To fire R at S under D , fire β .

We stipulate that an arbitrary rule R is equivalent, for given D and S , to a (D, S) -perspicuous rule R' obtained from R by renaming the bound variables. The equivalence means that $\text{Updates}(D, R, S) = \text{Updates}(D, R', S)$.

It remains to define $\beta = \text{Updates}(D, R, S)$ in the case when R is (D, S) -perspicuous.

If D is not empty, then β is the union of $\text{Updates}(\emptyset, R, S')$, where S' ranges over expansions of S such that $\text{Fun}(S') = \text{Fun}(S) \cup \text{Var}(D)$ and S' is consistent with D (so that the values of D variables are within their ranges in S'). Notice that $\beta = \emptyset$ if the range of any D variable is empty.

Suppose $D = \emptyset$. If R is an update instruction then $\beta = \text{Updates}(R, S)$. If R is a sequence of rules R_1, \dots, R_k , then β is the union of the update sets $\text{Updates}(\emptyset, R_i, S)$. Suppose that R is the conditional rule with clauses $(g_0, R_0), \dots, (g_k, R_k)$. Since R is covered by the empty declaration, the guards g_i have no free variables. We have two cases as usual. If all guards g_i fail in S , then β is empty, and if g_i is the first guard that holds in S then $\beta = \text{NUpdates}(\emptyset, R_i, S)$. Finally, if R is a declaration rule with declaration d and body R' then $\beta = \text{Updates}(d, R', S)$.

Remark Suppose that $D = \emptyset$ and R is a sequence of a declaration-free rule R_1 and a declaration rule R_2 with atomic declaration “Var v ranges over U ” followed by a declaration-free body R'_2 . Further suppose that U is empty in a state S appropriate for R and thus $\text{Updates}(D, R_2, S) = \emptyset$. Then $\text{Updates}(D, R, S)$ equals $\text{Updates}(D, R_1, S)$ which may be not empty. Contrary to the situation in 4.1.2, the empty range does not give inconsistency here. One cannot choose an element from the empty set, but one can execute a $R'_2(v)$ for every v in the empty set: just do not execute anything.

5.4 Importing Elements

The recursive definition of rules in 5.3 can be extended by import commands and/or (qualified) choice commands. The adjustment of the semantic definition is straightforward. For the sake of definiteness, consider the extension by means of the import constructor. The most important novelty, in comparison to 3.2.4, is that reserve elements have to be chosen for all combinations of the values of explicitly declared variables u such that the scope of the declaration of u properly includes the given import or choose subrule. For example, the rule

```
Var  $u$  ranges over  $U$ 
import  $v$ 
  Parent( $v$ ) :=  $u$ 
endimport
```

creates a new child for every element of U , and of course all these new children are different.

To reflect the novelty we redefine the domain of the global choice function. Suppose that D is a variable declaration, R is a rule covered by D , S is a state of an auxiliary vocabulary appropriate for R , and R is (D, S) -perspicuous. For every bound variable v of R , list all explicitly declared variables u such that either u occurs in D or u occurs in R and the scope of the declaration of u properly includes the scope of the declaration of v : u_1, \dots, u_l . (The adverb properly is there to exclude v from the list.) Let U_1, \dots, U_l be the ranges of u_1, \dots, u_l in S respectively, and \bar{U}_v be the Cartesian product $U_1 \times \dots \times U_l$. The desired global function ξ assigns different reserve elements to every pair (v, \bar{a}) where $v \in \text{Bound}(R)$ and $\bar{a} \in \bar{U}_v$.

Here is a variant of the example from 3.2.5:

```
Var  $u$  ranges over  $U$ 
extend Nodes with  $v_1, v_2$ 
  if Leaf( $u$ ) then
    FirstChild( $u$ ) :=  $v_1$ 
    SecondChild( $u$ ) :=  $v_2$ 
    NextSib( $v_1$ ) :=  $v_2$ 
  endif endextend
```

Remark Should one provide means to say explicitly that the main existential variable of a given choose rule depends only on such and such of the free variables of the rule? Maybe. But the need for such means has not been demonstrated yet.

5.5 Runs

Runs are defined as above.

6 Distributed Evolving Algebras

In this section we consider multi-agent computations. We do not suppose that agents are deterministic or do only a bounded amount of work at each step. The program of an agent may be any program described above.

Agents may share functions, and it is convenient [9] to assume that all states of all agents share the same carrier; see the end of 3.2.4 in this connection.

6.1 The Self Function

There is an interesting problem of self identification. It can be illustrated on the example of the following simple version of Dijkstra's dining philosophers protocol (which may deadlock). There are n philosophers, marked with numbers modulo n , each equipped with a fork. A philosopher i may think (which requires no forks) or eat using his/her fork and the fork of philosopher $i + 1$. A fork cannot be used by two philosophers at the same time.

Using functions

$$\text{Fork}_i = \begin{cases} \text{up} & \text{if the fork of philosopher } i \text{ is used,} \\ \text{down} & \text{otherwise,} \end{cases}$$

we can write a separate program for each philosopher i . Intuitively, however, all philosophers use the same program in the protocol.

To solve such problems, we suppose that each agent a is represented by an element of the common carrier. For simplicity, we will not distinguish between an agent and the element that represents the agent. Further, we use a special nullary function *Self*, interpreted differently by different agents. An agent a interprets *Self* as a . Thus function *Self* allows an agent to identify itself among other agents. *Self* is a logic name and cannot be the subject of an update instruction. To make rules sound a little better for humans, we use some capitalized pronouns, *e.g.* *Me*, as aliases for *Self*. Viewing agents as elements of the carrier is useful for other purposes as well. For example, it allows us to model the creation of new agents.

We return to the dining philosophers protocol. Here is a possible program (courtesy of Jim Huggins):

```
if Mode(Me)=think and Fork(Me)=Fork(Me+1)=down then
  Fork(Me):=up, Fork(Me+1):=up, Mode(Me):=eat
elseif Mode(Me)=eat then
  Fork(Me):=down, Fork(Me+1):=down, Mode(Me):=think
endif
```

It may be convenient to suppress the argument *Self*. For example, terms *Mode(Me)*, *Fork(Me)* and *Fork(Me+1)* may be treated as nullary functions and abbreviated, *e.g.*, as *mode*, *lfork* and

rfork, so that the rfork function of philosopher i is the lfork function of philosopher $i + 1$ and mode is a private function.

6.2 Basic Definition of Distributed Ealgebras

A *distributed ealgebra* \mathcal{A} consists of the following:

- A finite indexed set of single-agent programs π_ν , called *modules*. The *module names* ν are static nullary function names.
- A vocabulary $\Upsilon = \text{Fun}(\mathcal{A})$ which includes each $\text{Fun}(\pi_\nu) - \{\text{Self}\}$ but does not contain Self. In addition, Υ contains a unary function name Mod.
- A collection of Υ -states, called *initial states* of \mathcal{A} , satisfying the following conditions:
 - Different module names are interpreted as different elements.
 - There are only finitely many elements a such that, for some module name ν , $\text{Mod}(a) = \nu$.

A state S of vocabulary $\text{Fun}(\mathcal{A})$ is a state of \mathcal{A} if it satisfies the two conditions imposed on initial states. In applications it may make sense to restrict further the notion of state of the ealgebra in question.

An element a is an *agent* at S if there is a module name ν such that $S \models \text{Mod}(a) = \nu$; the corresponding π_ν is the program $\text{Prog}(a)$ of a , and $\text{Fun}(\pi_\nu)$ is the vocabulary $\text{Fun}(a)$ of a . Agent a is deterministic if $\text{Prog}(a)$ is so.

$\text{View}_a(S)$ is the reduct of S to vocabulary $\text{Fun}(a) - \{\text{Self}\}$ expanded with Self, which is interpreted as a . Think about $\text{View}_a(S)$ as the local state of agent a corresponding to the global state S . (It is not necessary to define local states via global states; see [8] for example.)

An agent a can *make a move* at S by firing $\text{Prog}(a)$ at $\text{View}_a(S)$ and changing S accordingly. As a part of the move, a may create new agents, *e.g.*, by importing reserve elements.

To perform a move of a deterministic agent a , fire

$$\text{Updates}(a, S) = \text{Updates}(\text{Prog}(a), \text{View}_a(S)).$$

Runs of a distributed ealgebra are defined below.

Cooperative Actions Consider a simple scenario with agents Sender and Receiver. If both are in mode Ready then Sender passes a value t_1 to Receiver who stores it at location $f(t_2)$. The transaction is atomic (that is, indivisible), but the Sender does not have access to $f(t_2)$ and the Receiver does not have access to t_1 , and thus neither agent is able to perform the transaction. A special auxiliary agent is needed to do the job, and it may be convenient to view the auxiliary agent as a team with members Sender and Receiver. Using functions Member₁ and Member₂ to specify the members of the team, we may write the following rule for the team, where Us is an alias for Self:

```

if Mode(Member1(Us))=Mode(Member2(Us))=Ready then
   $f(t_2) := t_1$ 
endif

```

In a similar way, one may have larger teams. Depending on need, teams may or may not be ordered.

6.3 Generalizations

6.3.1 Active Agents

Alter definition 6.2 as follows: Require that $\text{Fun}(\mathcal{A})$ contains an additional unary relation name *Active* and that only agents satisfying the relation *Active* (*active agents*) can make moves. This is essentially a generalization; the original definition can be seen as a special case where all agents are active.

The new definition may be convenient, for example, when the initial state specifies all agents and their programs, and these agents are activated and deactivated during the evolution.

The same convenience can be achieved without altering the original definition. (This may be useful, for example, if you want to prove something about all distributed ealgebras and wish to restrict attention to the basic definition without losing generality.) Here is one way to do that. In order to indicate the program of a potential agent without making it an actual agent, use an auxiliary unary function name *Mod'*. $\text{Active}(t)$ can be viewed as an abbreviation for $\text{Mod}(t) = \text{Mod}'(t)$ except if *Active* is the subject of an update instruction.

$\text{Active}(t) := t_0$

can be viewed as an abbreviation for

```
if  $t_0$  then  $\text{Mod}(t) := \text{Mod}'(t)$ 
else  $\text{Mod}(t) := \text{undef}$ 
endif
```

6.3.2 Active Teams

The generalized definition of distributed ealgebras described in 6.3.1 is used in this sub-subsection. The following problem was raised by Dean Rosenzweig [16].

Consider a scenario with (agents called) players and (additional agents viewed as) teams. Players form a static universe, teams form another static universe, and each agent is assigned a program once and for all. Players are activated and deactivated during the evolution. A team is supposed to be active if and only if its members are active. It follows that activating one player may necessitate the tedious work of activating many teams. Is there a simple and elegant way to ensure that every team is active when and only when all its members are active?

One obvious solution is to make teams active all the time and augment the program of each team with a guard stating that all the members are active. A more radical solution is to make the notion of team a part of the logic of distributed ealgebras. It will be ensured automatically that a team is active if and only if all its members are active. (It may be also required that the moves made by a player or any team involving the player are linearly ordered; see the second property of runs in 6.5.1 in this connection.) If substantial programming convenience is demonstrated, use that solution.

The possibilities to pay a lesser price in logic complication for the advantages of the radical solution will be discussed elsewhere. In this connection, Rosenzweig suggested generalizing further the definition of 6.3.1 by letting a possibly compound Boolean term play the role of *Active*. For example, $\text{Active}(v)$ may say that either v is a player satisfying an auxiliary relation *Ac* or v is a team with all members satisfying *Ac*.

6.4 Sequential Runs

We return to the basic definition of distributed ealgebras in 6.2.

A *pure sequential run* ρ of an ealgebra \mathcal{A} is a sequence $\langle S_n : n < \kappa \rangle$ of states of \mathcal{A} , where S_0 is an initial state and every S_{n+1} is obtained from S_n by executing a move of an agent. The generalization to the case of external functions or external locations is relatively straightforward.

Stages Since S_i may be equal to S_j for some $i \neq j$, it may be convenient to speak about stages. Starting from stage 0, the run goes through stages 1, 2, *etc.*. Formally, stage i can be defined as the pair (i, S_i) .

Quasi-sequential Runs An obvious generalization of a sequential run is a quasi-sequential run $\langle S_n : n < \kappa \rangle$, where each S_{n+1} is obtained from S_n by firing a collection A_n of agents. We do not mean that A_n is a team; since teams are agents, the definition of sequential runs does not exclude team moves. We mean that each $a \in A_n$ makes a move at S_n . If all agents are deterministic, then S_{n+1} is the result of firing $\bigcup \{\text{Updates}(S, a) : a \in A_n\}$.

Quasi-sequential runs may arise, for example, if you order moves in real (physical) time.

6.5 Partially Ordered Runs

Partially ordered computations are well known in the literature [12], [14], [11], *etc.* but we need to define our own version of that notion for our purposes here. We restrict attention to the case where moves are atomic and we use global states. Non-atomic moves have been explored in [3]. A simple notion of runs in [8] does not use global states.

Let us recall some well known notions. A *poset* is a partially ordered set. An *initial segment* of a poset P is a substructure X of P such that if $x \in X$ and $y < x$ in P then $y \in X$. Since X is a substructure, $y < x$ in X if and only if $y < x$ in P whenever $x, y \in X$. A *linearization* of a poset P is a linearly ordered set P' with the same elements such that if $x < y$ in P then $x < y$ in P' .

6.5.1 Runs

For simplicity, we restrict attention to pure runs and deterministic agents. A *run* ρ of a distributed ealgebra \mathcal{A} can be defined as a triple (M, A, σ) satisfying the following conditions 1–4.

- 1 M is a partially ordered set, where all sets $\{y : y \leq x\}$ are finite.

Elements of M represent *moves* made by various agents during the run. If $y < x$ then x starts when y is already finished; that explains why the set $\{y : y \leq x\}$ is finite.

- 2 A is a function on M such that every nonempty set $\{x : A(x) = a\}$ is linearly ordered.

$A(x)$ is the agent performing move x . The moves of any single agent are supposed to be linearly ordered.

- 3 σ assigns a state of \mathcal{A} to the empty set and each finite initial segment of M ; $\sigma(\emptyset)$ is an initial state.

$\sigma(X)$ is the result of performing all moves in X .

- 4 The coherence condition: If x is a maximal element in a finite initial segment X of M and $Y = X - \{x\}$, then $A(x)$ is an agent in $\sigma(Y)$ and $\sigma(X)$ is obtained from $\sigma(Y)$ by firing $A(x)$ at $\sigma(Y)$.

Intuitively, a run can be seen as the common part of histories of the same computation recorded by various observers. We hope to address this issue elsewhere.

If agents are not necessarily deterministic, we have to define moves as state transformers and make the coherence condition more precise:

- 4* If x is a maximal element in a finite initial segment X of M and $Y = X - \{x\}$, then $A(x)$ is an agent in $\sigma(Y)$, x is a move of $A(x)$ and $\sigma(X)$ is obtained from $\sigma(Y)$ by performing x at $\sigma(Y)$.

A run ρ' is an *initial segment* of a run ρ if (i) the move poset of ρ' is an initial segment of the move poset of ρ and (ii) the agent and state functions of ρ' are restrictions of those in ρ . A run ρ' is a *linearization* of ρ if the move poset of ρ' is a linearization of that of ρ , the agent function of ρ' is that of ρ , and the state function of ρ' is a restriction of that of ρ . Linearizations are sequential runs. A state S of is *reachable* in a run ρ if it belongs to the range of the state function of ρ .

Corollary 6.1 *All linearizations of the same finite initial segment of ρ have the same final state.*

Corollary 6.2 *A property holds in every reachable state of a run ρ if and only if it holds in every reachable state of every linearization of ρ .*

6.6 Real-time Computations

Real-time semantics appears in [3]. Ealgebras with clocks made their debut in [8]. We will have to address the issue of real time elsewhere.

Bibliography

References

- [1] Egon Börger, “Logic Programming: The Evolving Algebra Approach”, in PS.
- [2] Andreas Blass and Yuri Gurevich, “Evolving Algebras and Linear Time Hierarchy”, in PS.
- [3] Egon Börger, Yuri Gurevich and Dean Rosenzweig, “The Bakery Algorithm: Yet Another Specification and Verification”, this volume.
- [4] Egon Börger and Dean Rosenzweig, “The WAM — Definition and Compiler Correctness”, to appear in “Logic Programming: Formal Methods and Practical Applications”, Eds. C. Beierle and L. Pluemer, North-Holland, 1994.
- [5] Yuri Gurevich, “Logic and the challenge of computer science”, In “Current Trends in Theoretical Computer Science”, Ed. E. Börger, Computer Science Press, 1988, 1–57.
- [6] Yuri Gurevich, “Evolving Algebras: An Attempt to Discover Semantics ”, Bull. EATCS 43 (1991), 264–284; a slightly revised version in “Current Trends in Theoretical Computer Science”, Eds. G. Rozenberg and A. Salomaa, World Scientific, 1993, 266–292.

- [7] Yuri Gurevich and James K. Huggins, “The Semantics of the C Programming Language”, Springer Lecture Notes in Computer Science 702, 1993, 274–308.
- [8] Yuri Gurevich and Raghu Mani, “Group Membership Protocol: Formal Specification and Verification”, this volume.
- [9] Paola Glavan and Dean Rosenzweig, “Communicating Evolving Algebras”, in “Computer Science Logic”, eds. E. Börger et al., Lecture Notes in Computer Science 702, Springer, 1993, 182–215.
- [10] James K. Huggins, “Kermit: Specification and Verification”, this volume.
- [11] Shmuel Katz and Doron Peled, “Defining Conditional Independence Using Collapses”, Theoretical Computer Science 101 (1992), 337–359.
- [12] Leslie Lamport, “On Interprocess Communication”, Distributed Computing 1 (1986), 77–101.
- [13] Yuri Matijasevich, “Enumerable sets are Diophantine”, Soviet Math. Doklady 11:2 (1970), 354–358.
- [14] Antoni Mazurkiewicz, “Trace Theory”, Springer Lecture Notes in Computer Science 255 (1987), 279–324.
- [15] B. Pehrson and I. Simon, Editors, “IFIP 13th World Computer Congress 1994, Volume I: Technology/Foundations”, Elsevier, Amsterdam, to appear.
- [16] Dean Rosenzweig, Private Communication.